

Fahim Tajwar

Website: <https://tajwarfahim.github.io/> Email: tajwar93@stanford.edu

EDUCATION

Stanford University

Master of Science (MS), Computer Science (AI/ML)
Bachelor of Science (BS) with Distinction, Mathematics

Stanford, CA

June 2023 (Expected)
June 2022 (4.04/4.30)

PUBLICATIONS (* Equal Contribution)

Surgical Fine-Tuning Improves Adaptation to Distribution Shifts

2022

Yoonho Lee*, Annie S Chen*, Fahim Tajwar, Ananya Kumar, Huaxiu Yao, Percy Liang, Chelsea Finn
Under Review in International Conference on Learning Representations (ICLR), 2023

When to Ask for Help: Proactive Interventions in Autonomous Reinforcement Learning

2022

Annie Xie*, Fahim Tajwar*, Archit Sharma*, Chelsea Finn
Advances in Neural Information Processing Systems (NeurIPS), 2022
RSS Workshop on Scaling Robot Learning (SRL), 2022 (Spotlight)

Do Deep Networks Transfer Invariances Across Classes?

2022

Allan Zhou*, Fahim Tajwar*, Alexander Robey, Tom Knowles, George J. Pappas, Hamed Hassani, Chelsea Finn
International Conference on Learning Representations (ICLR), 2022

No True State-of-the-Art? OOD Detection Methods are Inconsistent across Datasets

2021

Fahim Tajwar, Ananya Kumar*, Sang Michael Xie*, Percy Liang
ICML Workshop on Uncertainty & Robustness in Deep Learning (UDL), 2021

Scalable deep learning to identify brick kilns and aid regulatory capacity

2021

Jihyeon Lee*, Nina R. Brooks*, Fahim Tajwar, Marshall Burke, Stefano Ermon, David B. Lobell, Debashish Biswas, Stephen P. Luby
Proceedings of the National Academy of Sciences, Apr 2021, 118 (17)

RESEARCH EXPERIENCE

Research Intern, Stanford Artificial Intelligence Laboratory

March 2020 – Current

- (Prof. Chelsea Finn and Percy Liang) Studied the relationship between distribution shifts and the layer of a neural network that should be fine-tuned (surgical fine-tuning) on the unsupervised (test-time) adaptation setting, and showed that adapting earlier layers instead of later/all layers help for corruption datasets like CIFAR-10-C and ImageNet-C (*Under review in ICLR, 2023*)
- (Prof. Chelsea Finn) Designed and implemented a reinforcement learning algorithm that works on environments with irreversibility (stuck states) under a variety of settings, e.g., episodic, autonomous, and continual learning (*NeurIPS, 2022*)
- (Prof. Chelsea Finn) Designed a generative model-based algorithm that can be combined with other state-of-the-art methods to give balanced test accuracy a boost of 1-2% on long-tailed versions of familiar datasets, e.g., CIFAR, GTSRB, etc. (*ICLR, 2022*)
- (Prof. Percy Liang) Worked on the Out-of-Distribution detection problem and showed different methods work better for different cases and there is no single “state-of-the-art” method, implying that the problem is too broad. (*ICML UDL Workshop, 2021*)

Computer Vision Research Intern, Stanford University

March 2019 – June 2020

- (Prof. Steve Luby and Stefano Ermon) Built a system using convolutional neural networks (CNNs), that detects environmental regulation violations in the form of brick kilns and produces their co-ordinates from satellite imagery in South Asia (*PNAS, 2021*)
- Detected nearly 10,000 brick kilns in Bangladesh which directly affect the lives of at least 1 million people using the trained CNN in *TensorFlow*, with the possibility of extending the project over to India (with 100,000 kilns) and other developing countries
- Created a classifier to distinguish between environment-friendly (Zigzag) and unfriendly (FCK) type of brick kilns

Research Intern, SLAC National Accelerator Laboratory

June 2018 – September 2018

- (Johanna Nelson Weker and Prof. Piero Pianetta) Using *Python* (“*Tomopy*”), generated 3D X-ray images of Lithium-Ion batteries from 2D projections and used the *Dragonfly* software to segment them
- Quantified the relationship between the morphological change near Anode (e.g., Lithium plating) and various accelerated ageing conditions --- saw a 40% increase in Lithium plating in batteries cycled at 55° C compared to those cycled at 25° C.

TEACHING EXPERIENCE

Teaching Assistant, CS 330 (Deep Multi-Task and Meta Learning), Stanford University	Sept 2022 – Dec 2022
Academic Tutor, Athletic Academic Resource Center (AARC), Stanford University	Sept 2021 – June 2022
Academic Tutor, Stanford University Mathematical Organization (SUMO)	Sept 2019 – June 2020
Trainer, National Physics Olympiad Team, Bangladesh (BdPhO)	Feb 2017 – June 2018

INDUSTRY EXPERIENCE

Software Engineer Intern, Meta Platforms (formerly Facebook) June 2022 – September 2022

- As part of the Ads Core ML Eng team, designed and implemented components for state-of-the-art ML recommendation systems in *Python*
- Experimented with various knowledge distillation techniques to improve performance of computationally cheaper ML networks
- Designed new modules with scalability in mind to make sure they work well with extremely large datasets and can also be trained efficiently

Software Engineer Intern, Cadence Design Systems June 2020 – September 2020

- Using *C++*, designed efficient data storage systems for graph neural networks (GNN), which reduced memory usage by 86%
- Using *Python* (*NumPy*), designed and implemented efficient data pre-processing modules
- Designed time and memory efficient data-loader classes in *Keras* and *TensorFlow*; researched memory-time tradeoff to find the optimal design

SKILLS

Programming Languages: C, C++, Python, Java, Matlab

Frameworks: Unix, PyTorch, Caffe2, TensorFlow

TALKS & PRESENTATION

- | | |
|--|-------------|
| • International Conference on Learning Representations (ICLR) | April 2022 |
| • ICML Workshop on Uncertainty & Robustness in Deep Learning (UDL) | July 2021 |
| • Stanford Earth Summer Undergraduate Research (SESUR) | August 2019 |
| • Stanford EE Research Experience for Undergraduates (REU) | August 2018 |

AWARDS

University Distinction, top 15% of the graduating class, Stanford University	June 2022
Bronze Medal, 48 th International Physics Olympiad, Indonesia	July 2017
Bronze Medal, 47 th International Physics Olympiad, Switzerland Liechtenstein	July 2016

SELECTED COURSEWORK

Mathematics: Real Analysis, Functional Analysis, Differential Topology, Measure Theory, Probability Theory, Graph Theory, Abstract Algebra, Linear Algebra & Matrix Theory, Statistical Inference, Numerical Computing

Computer Science: Machine Learning, Reinforcement Learning, Convex Optimization, Deep Learning for Computer Vision, Artificial Intelligence: Principles and Techniques