

Fahim Tajwar

Website: <https://tajwarfahim.github.io/> Email: ftajwar@cs.cmu.edu

EDUCATION

Carnegie Mellon University

Doctor of Philosophy (PhD), Machine Learning
Advisor: Ruslan Salakhutdinov & Jeff Schneider

Pittsburgh, PA

2023 – Current

Stanford University

Master of Science (MS), Computer Science (AI/ML)
Bachelor of Science (BS) with Distinction, Mathematics

Stanford, CA

2022 -- 2023

2017 -- 2022

PUBLICATIONS (* Equal Contribution)

Offline Retraining for Online RL: Decoupled Policy Learning to Mitigate Exploration Bias

2023

Max Sobol Mark*, Archit Sharma*, [Fahim Tajwar](#), Rafael Rafailov, Sergey Levine, Chelsea Finn

Under review, 2023

Conservative Prediction via Data-Driven Confidence Minimization

2023

Caroline Choi*, [Fahim Tajwar](#)*, Yoonho Lee*, Huaxiu Yao, Ananya Kumar, Chelsea Finn

Under review, 2023

Surgical Fine-Tuning Improves Adaptation to Distribution Shifts

2023

Yoonho Lee*, Annie S Chen*, [Fahim Tajwar](#), Ananya Kumar, Huaxiu Yao, Percy Liang, Chelsea Finn

International Conference on Learning Representations (ICLR), 2023

When to Ask for Help: Proactive Interventions in Autonomous Reinforcement Learning

2022

Annie Xie*, [Fahim Tajwar](#)*, Archit Sharma*, Chelsea Finn

Conference on Neural Information Processing Systems (NeurIPS), 2022

Do Deep Networks Transfer Invariances Across Classes?

2022

Allan Zhou*, [Fahim Tajwar](#)*, Alexander Robey, Tom Knowles, George J. Pappas, Hamed Hassani, Chelsea Finn

International Conference on Learning Representations (ICLR), 2022

No True State-of-the-Art? OOD Detection Methods are Inconsistent across Datasets

2021

[Fahim Tajwar](#), Ananya Kumar*, Sang Michael Xie*, Percy Liang

ICML Workshop on Uncertainty & Robustness in Deep Learning (UDL), 2021

Scalable deep learning to identify brick kilns and aid regulatory capacity

2021

Jihyeon Lee*, Nina R. Brooks*, [Fahim Tajwar](#), Marshall Burke, Stefano Ermon, David B. Lobell, Debashish Biswas, Stephen P. Luby

Proceedings of the National Academy of Sciences, Apr 2021, 118 (17)

RESEARCH EXPERIENCE

Research Intern, Stanford Artificial Intelligence Laboratory

March 2020 – June 2023

- (Prof. Chelsea Finn and Percy Liang) Relationship between distribution shifts and the layer of a neural network that should be fine-tuned (surgical fine-tuning) on the unsupervised (test-time) adaptation setting, specifically adapting earlier layers outperform adapting later/all layers for corruption datasets like CIFAR-10-C and ImageNet-C (*Under review in ICLR, 2023*)
- (Prof. Chelsea Finn) Deep reinforcement learning for irreversible environments with applications to episodic, autonomous, and continuous learning setups (*NeurIPS, 2022*)
- (Prof. Chelsea Finn) Generative model-based algorithm that produces performance boost of 1-2% when combined with other state-of-the-art methods on long-tailed versions of datasets like CIFAR, GTSRB, etc. (*ICLR, 2022*)
- (Prof. Percy Liang) Demonstration of out-of-distribution (OOD) detection problem being too broad by showing that many well-known methods don't perform consistently on a comprehensive suite of benchmark datasets (*ICML UDL Workshop, 2021*)

Computer Vision Research Intern, Stanford University

March 2019 – June 2020

- (Prof. Steve Luby and Stefano Ermon) Built a system using convolutional neural networks (CNNs), that detects environmental regulation violations in the form of brick kilns and produces their co-ordinates from satellite imagery in South Asia (*PNAS*, 2021)
- Detected nearly 10,000 brick kilns in Bangladesh which directly affect the lives of at least 1 million people using the trained CNN in *TensorFlow*, with the possibility of extending the project over to India (with 100,000 kilns) and other developing countries
- Created a classifier to distinguish between environment-friendly (Zigzag) and unfriendly (FCK) type of brick kilns

TEACHING EXPERIENCE

Teaching Assistant, Math 20 (Calculus), Stanford University

Jan 2023 – March 2023

Teaching Assistant, CS 330 (Deep Multi-Task and Meta Learning), Stanford University

Sept 2022 – Dec 2022

Academic Tutor, Athletic Academic Resource Center (AARC), Stanford University

Sept 2021 – June 2022

Academic Tutor, Stanford University Mathematical Organization (SUMO)

Sept 2019 – June 2020

INDUSTRY EXPERIENCE

Software Engineer Intern, Meta Platforms (formerly Facebook)

June 2022 – September 2022

- As part of the Ads Core ML Eng team, designed and implemented components for state-of-the-art ML recommendation systems in *Python*
- Experimented with various knowledge distillation techniques to improve performance of computationally cheaper ML networks
- Designed new modules with scalability in mind to make sure they work well with extremely large datasets and can also be trained efficiently

Software Engineering Intern, Cadence Design Systems

June 2020 – September 2020

- Using *C++*, designed efficient data storage systems for graph neural networks (GNN), which reduced memory usage by 86%
- Using *Python* (*NumPy*), designed and implemented efficient data pre-processing modules
- Designed time and memory efficient data-loader classes in *Keras* and *TensorFlow*; researched memory-time tradeoff to find the optimal design

SKILLS

Programming Languages: Python, C, C++, Java, Matlab

Frameworks: Unix, PyTorch, Caffe2, TensorFlow

TALKS & PRESENTATION

- Neural Information Processing Systems (NeurIPS) November 2022
- International Conference on Learning Representations (ICLR) April 2022
- ICML Workshop on Uncertainty & Robustness in Deep Learning (UDL) July 2021
- Stanford Earth Summer Undergraduate Research (SESUR) August 2019
- Stanford EE Research Experience for Undergraduates (REU) August 2018

AWARDS

- Top Reviewer, Conference on Neural Information Processing Systems (NeurIPS) 2023
- University Distinction, top 15% of the graduating class, Stanford University 2022
- Tau Beta Pi Engineering Honor Society 2020
- Bronze Medal, 48th International Physics Olympiad, Indonesia 2017
- Bronze Medal, 47th International Physics Olympiad, Switzerland Liechtenstein 2016

SERVICE

- Reviewer, Conference on Neural Information Processing Systems (NeurIPS) (**Top Reviewer, 2023**) 2023
- Reviewer, NeurIPS Workshop on Distribution Shifts (DistShift) 2023
- Reviewer, International Conference on Learning Representations (ICLR) 2024
- Reviewer, International Conference on Machine Learning (ICML) 2024