

Fahim Tajwar

Website: <https://tajwarfahim.github.io/> Email: ftajwar@cs.cmu.edu

EDUCATION

Carnegie Mellon University

Doctor of Philosophy (PhD), Machine Learning
Advisor: Ruslan Salakhutdinov & Jeff Schneider

Pittsburgh, PA

2023 – Current

Stanford University

Master of Science (MS), Computer Science (AI/ML)
Bachelor of Science (BS) with Distinction, Mathematics

Stanford, CA

2023

2022

SELECTED PUBLICATIONS (* Equal Contribution)

Training a Generally Curious Agent

2025

[Fahim Tajwar*](#), Yiding Jiang*, Abitha Thankaraj, Sumaita Sadia Rahman, J Zico Kolter, Jeff Schneider, Ruslan Salakhutdinov
International Conference on Machine Learning (ICML), 2025 (**Spotlight**)

Preference Fine-Tuning of LLMs Should Leverage Suboptimal, On-Policy Data

2024

[Fahim Tajwar*](#), Anikait Singh*, Archit Sharma, Rafael Rafailov, Jeff Schneider, Tengyang Xie, Stefano Ermon, Chelsea Finn, Aviral Kumar
International Conference on Machine Learning (ICML), 2024

Conservative Prediction via Data-Driven Confidence Minimization

2024

Caroline Choi*, [Fahim Tajwar*](#), Yoonho Lee*, Huaxiu Yao, Ananya Kumar, Chelsea Finn
Transactions on Machine Learning Research (TMLR), 2024

Surgical Fine-Tuning Improves Adaptation to Distribution Shifts

2023

Yoonho Lee*, Annie S. Chen*, [Fahim Tajwar](#), Ananya Kumar, Huaxiu Yao, Percy Liang, Chelsea Finn
International Conference on Machine Learning (ICML), 2024

When to Ask for Help: Proactive Interventions in Autonomous Reinforcement Learning

2022

Annie Xie*, [Fahim Tajwar*](#), Archit Sharma*, Chelsea Finn
Conference on Neural Information Processing Systems (NeurIPS), 2022

Do Deep Networks Transfer Invariances Across Classes?

2022

Allan Zhou*, [Fahim Tajwar*](#), Alexander Robey, Tom Knowles, George J. Pappas, Hamed Hassani, Chelsea Finn
International Conference on Learning Representations (ICLR), 2022

Scalable Deep Learning to Identify Brick Kilns and Aid Regulatory Capacity

2021

Jihyeon Lee*, Nina R. Brooks*, [Fahim Tajwar](#), Marshall Burke, Stefano Ermon, David B. Lobell, Debashish Biswas, Stephen P. Luby
Proceedings of the National Academy of Sciences (PNAS), 2021

INDUSTRY EXPERIENCE

Software Engineer Intern, Meta Platforms

June 2022 – September 2022

Software Engineer Intern, Cadence Design Systems

June 2020 – September 2020

TALKS & PRESENTATION

- Neural Information Processing Systems (NeurIPS) November 2022
- International Conference on Learning Representations (ICLR) April 2022
- ICML Workshop on Uncertainty & Robustness in Deep Learning (UDL) July 2021
- Stanford Earth Summer Undergraduate Research (SESUR) August 2019
- Stanford EE Research Experience for Undergraduates (REU) August 2018

AWARDS

Top Reviewer, Conference on Neural Information Processing Systems (NeurIPS)	2023
University Distinction, top 15% of the graduating class, Stanford University	2022
Tau Beta Pi Engineering Honor Society	2020
Bronze Medal, 48 th International Physics Olympiad (IPhO), Indonesia	2017
Bronze Medal, 47 th International Physics Olympiad (IPhO), Switzerland Liechtenstein	2016

SERVICE

• Reviewer, Conference on Neural Information Processing Systems (NeurIPS) (Top Reviewer, 2023)	2023, 2024, 2025
• Reviewer, International Conference on Learning Representations (ICLR)	2024, 2025
• Reviewer, International Conference on Machine Learning (ICML)	2024, 2025
• Reviewer, The IEEE/CVF Computer Vision and Pattern Recognition Conference (CVPR)	2024
• Reviewer, Transactions on Machine Learning Research (TMLR)	2024-2025
• Reviewer, International Joint Conference on Artificial Intelligence (IJCAI)	2024